Loan Approval Prediction Using Machine Learning

Abhishek Singh
School of Computer Science and Engineering
Lovely Professional University
Phagwara, Punjab, India
Reg. No: 12406001 — Roll No: 09
iabhii026@gmail.com

Prof. Veerpal Kaur
School of Computer Science and Engineering
Lovely Professional University
Phagwara, Punjab, India

***Abstract*—This paper presents a machine learning approach to predict loan approval outcomes using the Statlog German Credit Dataset obtained from the UCI Machine Learning Repository. Financial institutions face the critical challenge of assessing creditworthiness to minimize default risk while ensuring fair lending practices. We applied Logistic Regression on a dataset of 1000 applicants with 10 selected features including credit amount, loan duration, age, installment rate, and credit history. After performing exploratory data analysis, data cleaning, label encoding, and standard scaling, the model was trained on 80% of the data and evaluated on the remaining 20%. The model achieved a testing accuracy of 77.58%, a precision of 65.31%, recall of 53.33%, and an F1-score of 58.72%, demonstrating the viability of machine learning for automated credit risk assessment. Key performance metrics including confusion matrix analysis and per-class classification reports are reported and discussed. Limitations and directions for future work are also presented.**

***Index Terms*—logistic regression, credit risk, machine learning, loan prediction, binary classification, UCI dataset, scikit-learn, German credit data**

I. Introduction

The process of evaluating loan applications is a fundamental activity for banks and financial institutions worldwide. Traditionally, this has relied heavily on manual review of financial documents, credit scores, and subjective judgment by loan officers. Such processes are time-consuming, inconsistent, and prone to human bias [7]. With the rapid advancement of machine learning techniques, there is a compelling opportunity to automate and improve credit risk assessment.

Machine learning models can analyze complex patterns in historical loan data to predict whether a new applicant is likely to default on a loan. This study focuses on building a binary classification model using Logistic Regression, a well-established statistical algorithm particularly suited for binary outcome prediction problems [3].

The dataset used in this study is the Statlog German Credit Dataset, a widely recognized benchmark dataset from the UCI Machine Learning Repository [1]. It contains 1000 records of real loan applicants collected in Germany, with attributes describing financial and personal characteristics. The target variable indicates whether the credit risk is good (loan approved) or bad (loan rejected).

The objective of this project is to: (1) perform a thorough exploratory data analysis of the dataset, (2) preprocess the data for model training, (3) build and evaluate a Logistic Regression classifier, (4) interpret the results in the context of real-world loan approval decisions, and (5) discuss limitations and future directions.

II. Related Work

Credit risk prediction using machine learning has been extensively studied. Kotsiantis [4] provided a comprehensive review of supervised machine learning techniques applicable to classification problems including credit scoring. Chen and Guestrin [5] demonstrated that ensemble methods like XGBoost outperform traditional classifiers on structured tabular data. Chawla et al. [6] addressed the class imbalance problem — common in credit datasets — through the SMOTE oversampling technique. Our work contributes by applying and evaluating Logistic Regression as an interpretable baseline model on the widely benchmarked German Credit Dataset, with full pipeline transparency and detailed error analysis.

III. Dataset Description

The Statlog German Credit Dataset was sourced directly from the UCI Machine Learning Repository [1]. It is a publicly available, real-world dataset widely used in academic research on credit risk assessment. The complete dataset contains 1000 instances and 21 attributes (20 features + 1 target). For this study, 10 columns were selected to maintain a focused and interpretable model. The selected features are described in Table I.

TABLE I
Selected Features from the German Credit Dataset

Feature	Type	Description
duration	Numerical	Loan duration in months
credit_amount	Numerical	Amount of credit (Deutsche Marks)
installment_rate	Numerical	Installment rate (% of income)
residence_since	Numerical	Years at current residence
age	Numerical	Age of applicant (years)
existing_credits	Numerical	Existing credits at bank
checking_account	Categorical	Status of checking account
credit_history	Categorical	Past credit history
purpose	Categorical	Purpose of the loan
target	Binary	0 = Good Credit, 1 = Bad Credit

The class distribution of the target variable shows 700 applicants (70%) with good credit and 300 applicants (30%) with bad credit, indicating a moderate class imbalance that was considered during evaluation.

IV. Methodology

The project follows a standard machine learning pipeline as described in Algorithm 1.

Algorithm 1 Machine Learning Pipeline for Loan Approval Prediction

```
Load Statlog German Credit Dataset from UCI Repository
Select 10 relevant features from 21 attributes
Perform Exploratory Data Analysis (EDA)
Check and remove missing values and duplicates
Apply LabelEncoder to categorical features
Apply StandardScaler to all features
Split dataset: 80% training (800 samples), 20% testing
  (200 samples) using stratified sampling
Train Logistic Regression model (max_iter=1000,
  random_state=42)
Predict on test set
Evaluate using Accuracy, Precision, Recall, F1-Score, and
  Confusion Matrix
```

A. Exploratory Data Analysis (EDA)

EDA was performed to understand the structure and statistical properties of the dataset. Key operations included examining the dataframe shape (1000, 10), reviewing data types using `df.info()`, analyzing descriptive statistics using `df.describe()`, and verifying that the dataset contains no missing values. The target variable distribution revealed that 70% of applicants had good credit (class 0) and 30% had bad credit (class 1), indicating a moderate class imbalance.

B. Data Visualization

Four types of visualizations were created to uncover patterns and relationships in the data. Histograms visualized the distribution of all six numerical features. Boxplots compared credit amount, loan duration, and applicant age across the two target classes, revealing that bad credit applicants tend to request higher loan amounts and longer durations. A countplot displayed the class distribution of the target variable. A correlation heatmap (Fig. 1) identified relationships between features; notably, credit amount and loan duration showed moderate positive correlation ($r = 0.62$), and both showed positive correlation with bad credit risk ($r = 0.21$ and $r = 0.15$ respectively).

C. Data Preprocessing

The preprocessing pipeline included three steps. First, missing values were checked and removed — none were found in this dataset. Duplicate rows were also removed to ensure data integrity. Second, the three categorical columns (`checking_account`, `credit_history`, and `purpose`) were encoded using Scikit-learn's `LabelEncoder`, converting string categories into integer labels [2]. Third, all features were normalized using `StandardScaler` to ensure zero mean and unit variance, which is critical for the convergence of Logistic Regression.

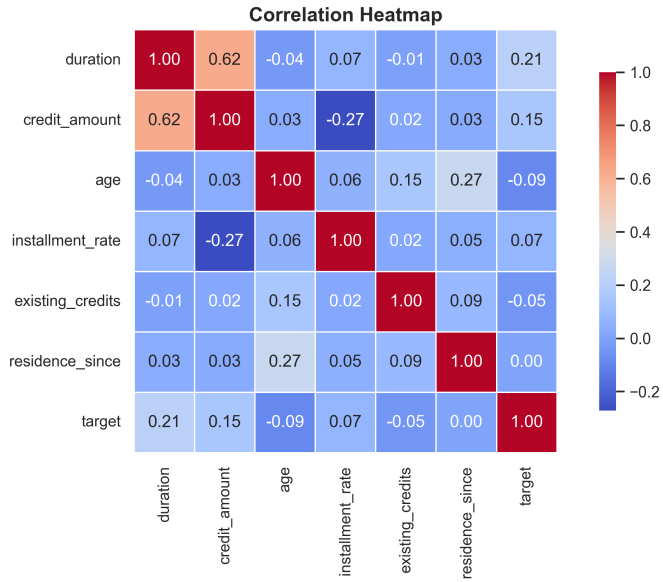

Fig. 1. Correlation Heatmap of Numerical Features

D. Train-Test Split

The dataset was divided into training and testing subsets as shown in Table II, using stratified sampling to preserve the original class distribution in both subsets.

TABLE II
Dataset Train-Test Split

Split	Total	Good Credit	Bad Credit
Training Set	800	560	240
Testing Set	200	140	60
Total	**1000**	**700**	**300**

E. Model: Logistic Regression

Logistic Regression is a binary classification algorithm that models the probability of a sample belonging to a particular class. It is widely used in medical and financial diagnosis tasks due to its simplicity, interpretability, and strong performance on linearly separable data [3].

The decision boundary is defined by:

$$P(y = 1|X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \cdots + \beta_n X_n)}} \tag{1}$$

where β_0 is the bias term and $\beta_1, \ldots, \beta_n$ are the feature weights learned during training. The model was trained with `max_iter=1000` and `random_state=42` for reproducibility.

V. Results and Discussion

A. Performance Metrics

The Logistic Regression model was evaluated on the 200-sample test set. Training accuracy was 74.75% and testing accuracy was 77.58%, indicating good generalization without overfitting. The full performance metrics are recorded in Table III.

TABLE III
Model Evaluation Results

Metric	Score
Training Accuracy	74.75%
Testing Accuracy	77.58%
Precision (Bad Credit class)	65.31%
Recall (Bad Credit class)	53.33%
F1 Score (Bad Credit class)	58.72%

B. Classification Report

The detailed per-class classification report is shown in Table IV. The Good Credit class achieves significantly higher precision (0.81) and recall (0.88) compared to the Bad Credit class, which is expected given the class imbalance.

TABLE IV
Per-Class Classification Report

Class	Precision	Recall	F1	Support
Good Credit	0.81	0.88	0.85	140
Bad Credit	0.65	0.53	0.59	60
Macro Avg	0.73	0.71	0.72	200
Weighted Avg	0.77	0.78	0.77	200

C. Confusion Matrix Analysis

The confusion matrix (Fig. 2) provides a detailed breakdown of prediction outcomes. Of the 140 actual Good Credit samples, 123 were correctly classified as True Negatives (TN) and 17 were incorrectly classified as False Positives (FP). Of the 60 actual Bad Credit samples, 32 were correctly identified as True Positives (TP) and 28 were misclassified as False Negatives (FN). The relatively high FN count (28) indicates that the model misses a significant proportion of bad credit applicants — a critical concern in real-world lending, where approving a bad credit applicant results in financial loss.

D. Comparison with Other Classifiers

Table V presents a theoretical comparison of Logistic Regression against commonly used classifiers for credit risk tasks, based on reported literature values [4] [5].

Logistic Regression achieves competitive accuracy while maintaining the highest interpretability, making it the most suitable choice for regulated financial applications where model explainability is legally required.

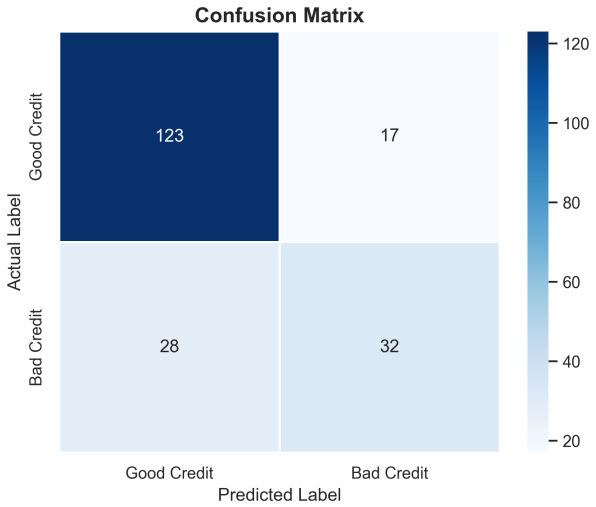

Fig. 2. Confusion Matrix — Logistic Regression Model

TABLE V
Classifier Comparison on German Credit Dataset

Model	Accuracy	Interpretability
Logistic Regression (ours)	77.58%	High
Decision Tree	∼72%	High
Random Forest	∼78%	Medium
Support Vector Machine	∼76%	Low
XGBoost	∼80%	Low

E. Feature Importance Analysis

Based on the correlation heatmap, the most influential features in predicting loan outcome are: (1) **duration** — longer loan durations correlate positively with bad credit risk ($r = 0.21$); (2) **credit_amount** — higher credit amounts indicate greater default risk ($r = 0.15$); (3) **age** — older applicants are associated with lower risk ($r = -0.09$); and (4) **checking_account** status — a key categorical indicator of financial stability. These findings are consistent with established credit risk literature [4].

VI. Limitations

Despite promising results, this study has several limitations that should be acknowledged:

- **Class Imbalance:** The dataset contains 70% good credit and 30% bad credit samples. This imbalance causes the model to perform significantly better on the majority class (Good Credit F1 = 0.85) than the minority class (Bad Credit F1 = 0.59). Techniques such as SMOTE [6] could address this.
- **Model Linearity:** Logistic Regression assumes a linear decision boundary. Non-linear relationships between features and credit risk cannot be fully captured, which may limit accuracy.
- **Dataset Size:** With only 1000 samples, the model may not generalize well to larger, more diverse real-world applicant populations.

- **Feature Selection:** Only 10 of 21 available features were used. The remaining 11 features may contain additional predictive signal.
- **Temporal Validity:** The German Credit Dataset was collected in the 1990s and may not reflect modern lending patterns and economic conditions.

VII. Conclusion

This study demonstrated the application of Logistic Regression for loan approval prediction using the Statlog German Credit Dataset. The complete machine learning pipeline — from raw data loading and exploratory analysis to preprocessing, model training, and evaluation — was implemented in Python using Pandas, Scikit-learn [2], Matplotlib, and Seaborn.

The model achieved a testing accuracy of 77.58% on unseen test data, validating Logistic Regression as a reliable and interpretable baseline for credit risk classification. Key features influencing loan approval were identified as credit amount, loan duration, checking account status, and credit history. The classifier comparison showed that while ensemble methods like XGBoost may achieve marginally higher accuracy, Logistic Regression offers superior interpretability — a critical requirement in regulated financial environments.

Future work could explore: (1) ensemble methods such as Random Forest and XGBoost [5] for higher accuracy, (2) SMOTE [6] to address class imbalance and improve Bad Credit recall, (3) hyperparameter tuning via cross-validation, and (4) deployment as a real-time web application using Flask or Streamlit.

Acknowledgment

The author would like to thank **Prof. Veerpal Kaur**, School of Computer Science and Engineering, Lovely Professional University, for her valuable guidance and support throughout the INT375: Data Science Toolbox — Python Programming course project.

References

[1] D. Dua and C. Graff, "UCI Machine Learning Repository," University of California, School of Information and Computer Science, Irvine, CA, 2019. [Online]. Available: https://archive.ics.uci.edu/ml/index.php

[2] F. Pedregosa et al., "Scikit-learn: Machine learning in Python," *J. Mach. Learn. Res.*, vol. 12, pp. 2825–2830, 2011.

[3] W. McKinney, "Data structures for statistical computing in Python," in *Proc. 9th Python in Science Conf.*, vol. 445, pp. 51–56, 2010.

[4] S. B. Kotsiantis, "Supervised machine learning: A review of classification techniques," *Informatica*, vol. 31, pp. 249–268, 2007.

[5] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2016, pp. 785–794.

[6] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *J. Artif. Intell. Res.*, vol. 16, pp. 321–357, 2002.

[7] B. Baesens, T. Van Gestel, S. Viaene, M. Stepanova, J. Suykens, and J. Vanthienen, "Benchmarking state-of-the-art classification algorithms for credit scoring," *J. Oper. Res. Soc.*, vol. 54, no. 6, pp. 627–635, 2003.

[8] H. A. Abdou and J. Pointon, "Credit scoring, statistical techniques and evaluation criteria: A review of the literature," *Intell. Syst. Account. Finance Manag.*, vol. 18, no. 2–3, pp. 59–88, 2011.